

The background features abstract green geometric shapes. On the left, a tall, narrow green triangle points downwards. On the right, a large, complex shape composed of several overlapping green triangles of varying shades (from light lime to dark forest green) is positioned. A thin, light gray line extends from the bottom left towards the right, passing behind the green shapes.

Container Virtualization

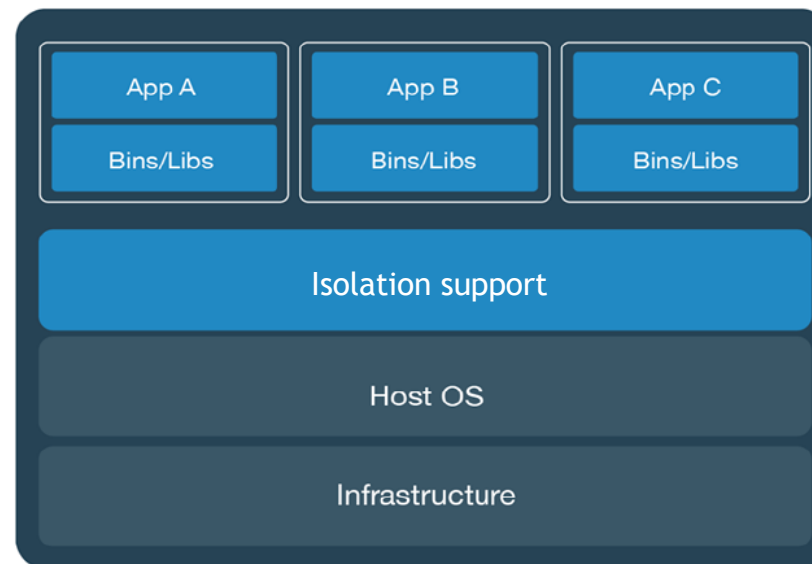
Diogo Gomes

OS Level Virtual Machines

- ▶ Do not provide a full virtualized hardware
 - ▶ Only OS services are virtualized
 - ▶ Host kernel: virtualize its services in a isolated manner to different guests
- ▶ Container: isolated execution environment to encapsulate one or more processes
- ▶ Relevant aspects:
 - ▶ Isolation/Visibility: limit what can be seen by tenants
 - ▶ Resource control: limit resource consumption
 - ▶ Portability: reconstruct the same environment in multiple hosts

Container based virtualization

- ▶ Container based virtualization relies on OS mechanisms for enhanced isolation



OS Support

- ▶ Linux has no support for containers!
 - ▶ That is: the Linux OS does not know what is a container
- ▶ Containers are a set of configurations, creating an environment, applied to a process (and its children)
 - ▶ Namespaces: isolation and virtualization of each tenant
 - ▶ Cgroups: resource control of CPU, RAM, IO of each tenant

Underlying technologies

- ▶ Namespaces
 - ▶ (mnt, pid, net, ipc, uts/hostname, user ids)
- ▶ cgroups
 - ▶ (cpu, memory, disk, i/o - resource management)
- ▶ AppArmor, SELinux
 - ▶ (security/access control)
- ▶ seccomp
 - ▶ (computation isolation)
- ▶ chroot
 - ▶ (file system isolation)

Namespaces

- ▶ Mechanism used to isolate and virtualize system resources
 - ▶ Processes in a namespace are unable to see remaining resources
 - ▶ Their namespace looks like the entire host
 - ▶ Processes in a namespace access resources as if there is no concurrency
- ▶ Network namespace:
 - ▶ Mechanism to create a somewhat independent set of networking resources
 - ▶ Network interfaces, routing tables
 - ▶ Network interfaces can only belong to a single namespace
- ▶ PID namespace:
 - ▶ Restricted namespace with private process identifiers
 - ▶ Host processes are invisible

cgroups

- ▶ Mechanism used to restrict (limit, control) or monitor the amount of resources used by “groups of processes”
 - ▶ Processes can be organized in groups, to control their accesses to resources
- ▶ Example: CPU control groups for scheduling
 - ▶ Limit the amount of CPU time that processes can use, etc...
- ▶ Similar cgroups for other resources
 - ▶ memory, IO, pids, network, ..

Setting up a container

1. Setup all the needed namespaces and control groups
 - ▶ Usually according to some template
2. Create a “disk image” for the container
 - ▶ directory containing the container’s filesystem
 - ▶ will exist in the host (it’s a directory)
3. Chroot to the container filesystem
 - ▶ Must contain all the libraries/files needed to execute the program
4. Start the program to containerize
 - ▶ This process it will have PID 1 in the container
 - ▶ Note: this process can mount procfs or other pseudo-file systems
 - ▶ Namespaces allow to control the information exported in those pseudofilesystems

Setting up a container

- ▶ Due to network namespace, containerized processes do not see the host's network interfaces
 - ▶ But we usually do not assign a real network interface to a namespace
 - ▶ Would have impact to the host
- ▶ Networking with containers:
 - ▶ Create a virtual ethernet pair: Two virtual interfaces, connected point-to-point
 - ▶ Packets sent on one interface are received on the other, and vice-versa
 - ▶ Associate one virtual interfaces to the network namespace of the container
 - ▶ Bind the other one to a software bridge

Containers are not Hardware VMs

- Core idea: a container is a **Process in a sandbox**, not a fully emulated virtual hardware

| Virtual Machine | Container |
|---|---|
| Each VM runs a different OS | All containers share the same OS |
| Booting the VM involves the Full boot of the guest HW | Containers boot in seconds (start service) |
| VM snapshots are made on demand | Containers use multiple images in layers |
| VM description is a XML document | Container description is a sequence of instructions |
| VM will pre-consume RAM and CPU | Containers have almost zero overhead |
| A VM description will result in one VM | A dockerfile may be used to start many containers |
| Has limits of tens of VMs per CPU core | Can have hundreds of containers per CPU core |

Tools

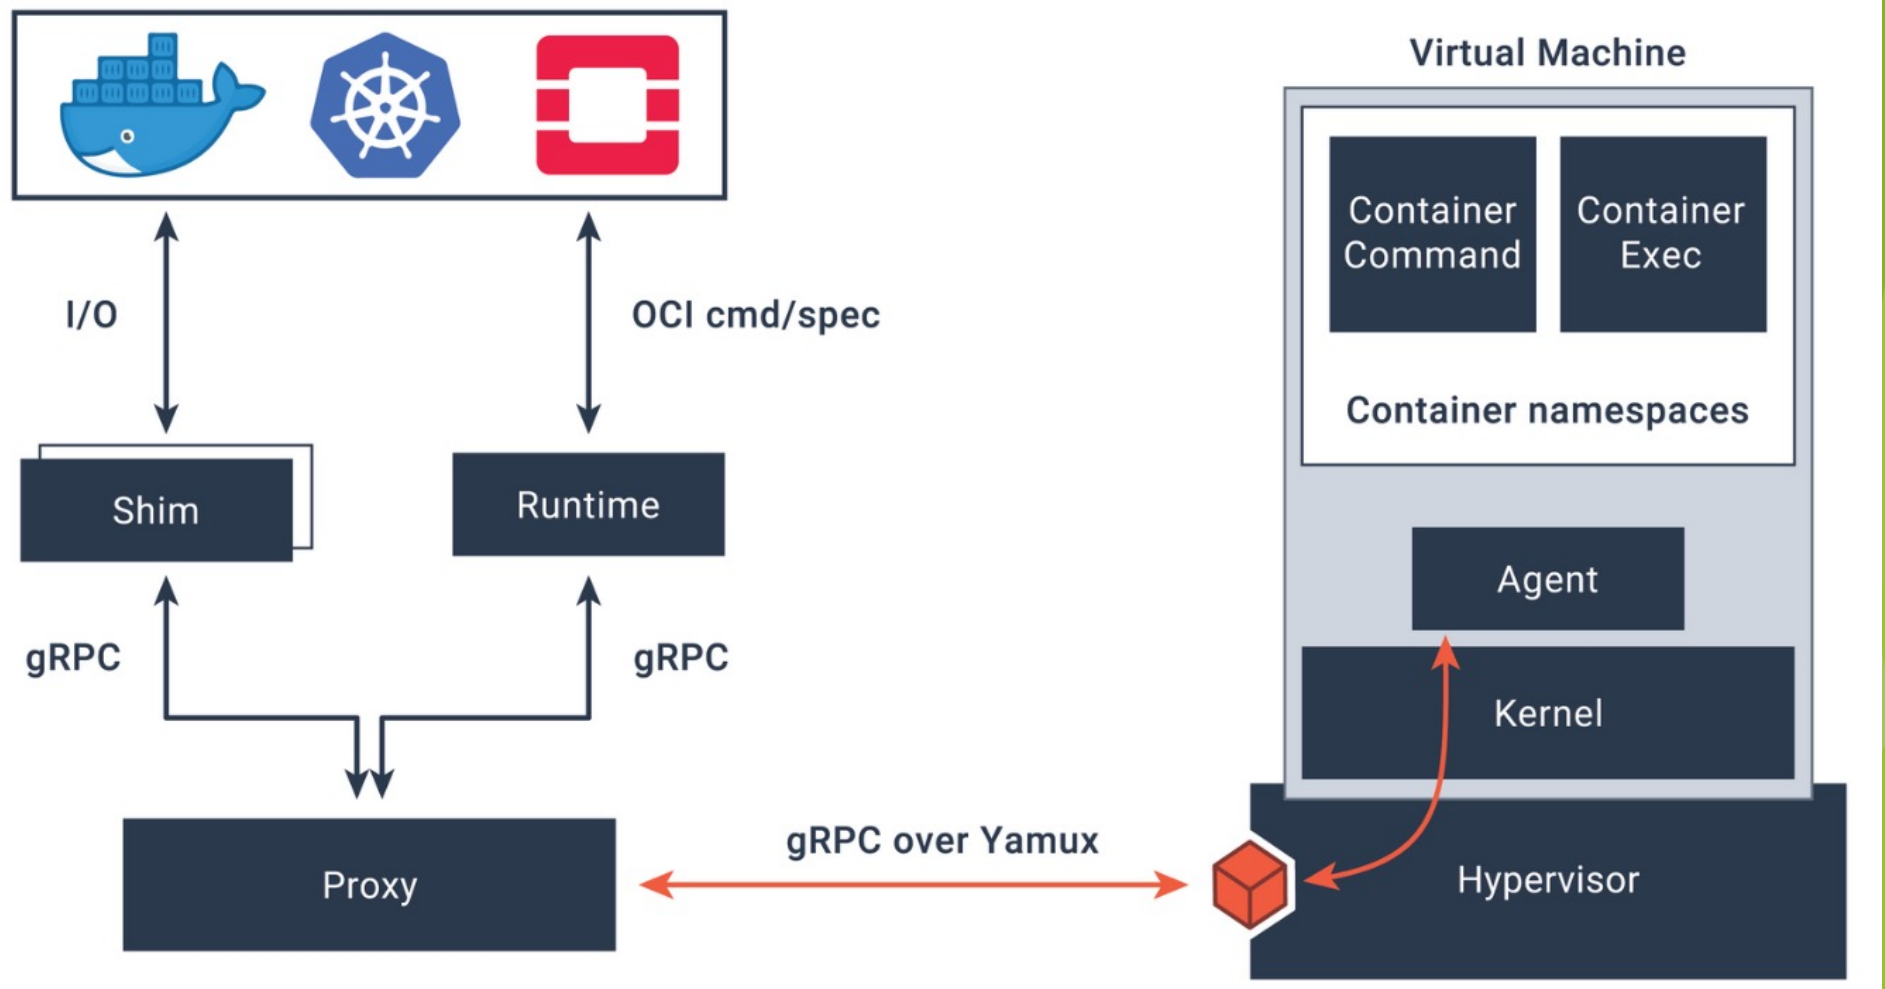
- ▶ Creating containers by hand is a difficult tasks
 - ▶ Reasonably high set of instructions to create namespaces, cgroups, chroots, interfaces, bridges
- ▶ User space tools facilitate manipulation of containers
 - ▶ LXC: Linux Containers
 - ▶ Docker
 - ▶ Singularity
 - ▶ Kubernetes

Portability

- ▶ Containers are all about creating replicable execution environments
 - ▶ Same environment (namespaces, chroot, cgroups) that can be used in different hosts
- ▶ Important to have a common description of containers
- ▶ Open Container Initiative (OCI): <https://www.opencontainers.org/>
 - ▶ Defines standards for user-space tools
 - ▶ Runtime specification: configuration, execution, and lifecycle of a container
 - ▶ Image specification: how to represent the data of a container


Portability

- ▶ Containers are evolving beyond OS level virtualization
 - ▶ If runtime and images are standardized, implementation doesn't matter.
 - ▶ It is not relevant if the container will run on OS Level Virtualization or on a HW VM
- ▶ Containers are evolving to Light Weight Virtual Machines
 - ▶ Namespaces: can be implemented as a Virtual Machines
 - ▶ Cgroups: can be implemented as the Virtual Machines with control tools
 - ▶ Images: can be implemented as virtual disks
 - ▶ Example: Kata containers



Docker

- ▶ Commercial Product
 - ▶ Evolution from Solaris Containers and Linux Containers (LXC)
 - ▶ At a conceptual level
 - ▶ Released as Open Source in 2013
- ▶ Core Concept: Container
 - ▶ App level construct (unrelated to infrastructure)
 - ▶ Composed by multiple functionality (namespaces) in coordination
 - ▶ One Container runs One Application (which can spawn child apps)
- ▶ Containers pack an application to run in any underlying hardware
 - ▶ including configurations, dependencies, auxiliary data, etc...

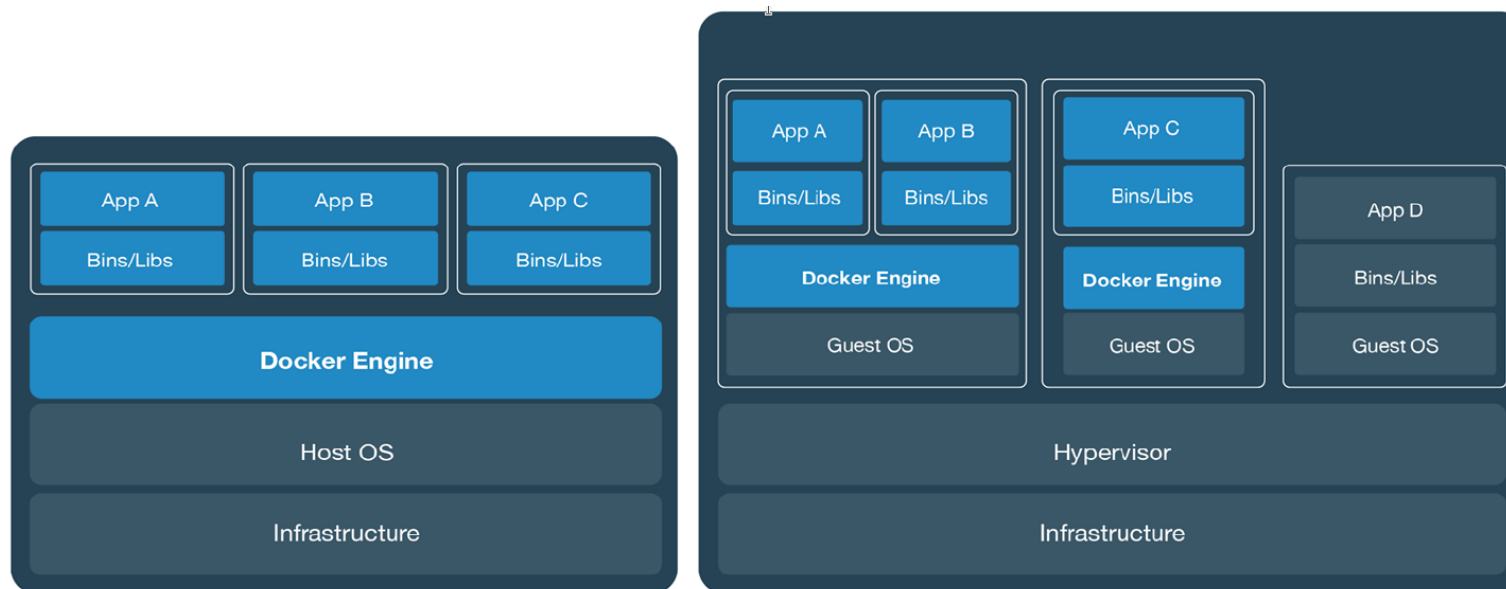


Docker Concepts

- ▶ **Image:** **the data** of a container (application, libraries, images, etc...)
- ▶ **Container:** A **running instance** of an application.
 - ▶ Composed by one or more images. Each image adds a specific functionality
- ▶ **Engine:** Software the executes commands for containers
- ▶ **Registry:** **Repository** of docker images
- ▶ **Control Plane:** infrastructure to manage containers and images



Containers can run with VMs





Linux Containers



liblxc

namespaces

cgroups

SELinux/AppArmor

Linux kernel



Docker 1.10 and later



runC

runC

runC

containerd-shim

containerd-shim

containerd-shim

containerd

Docker Engine

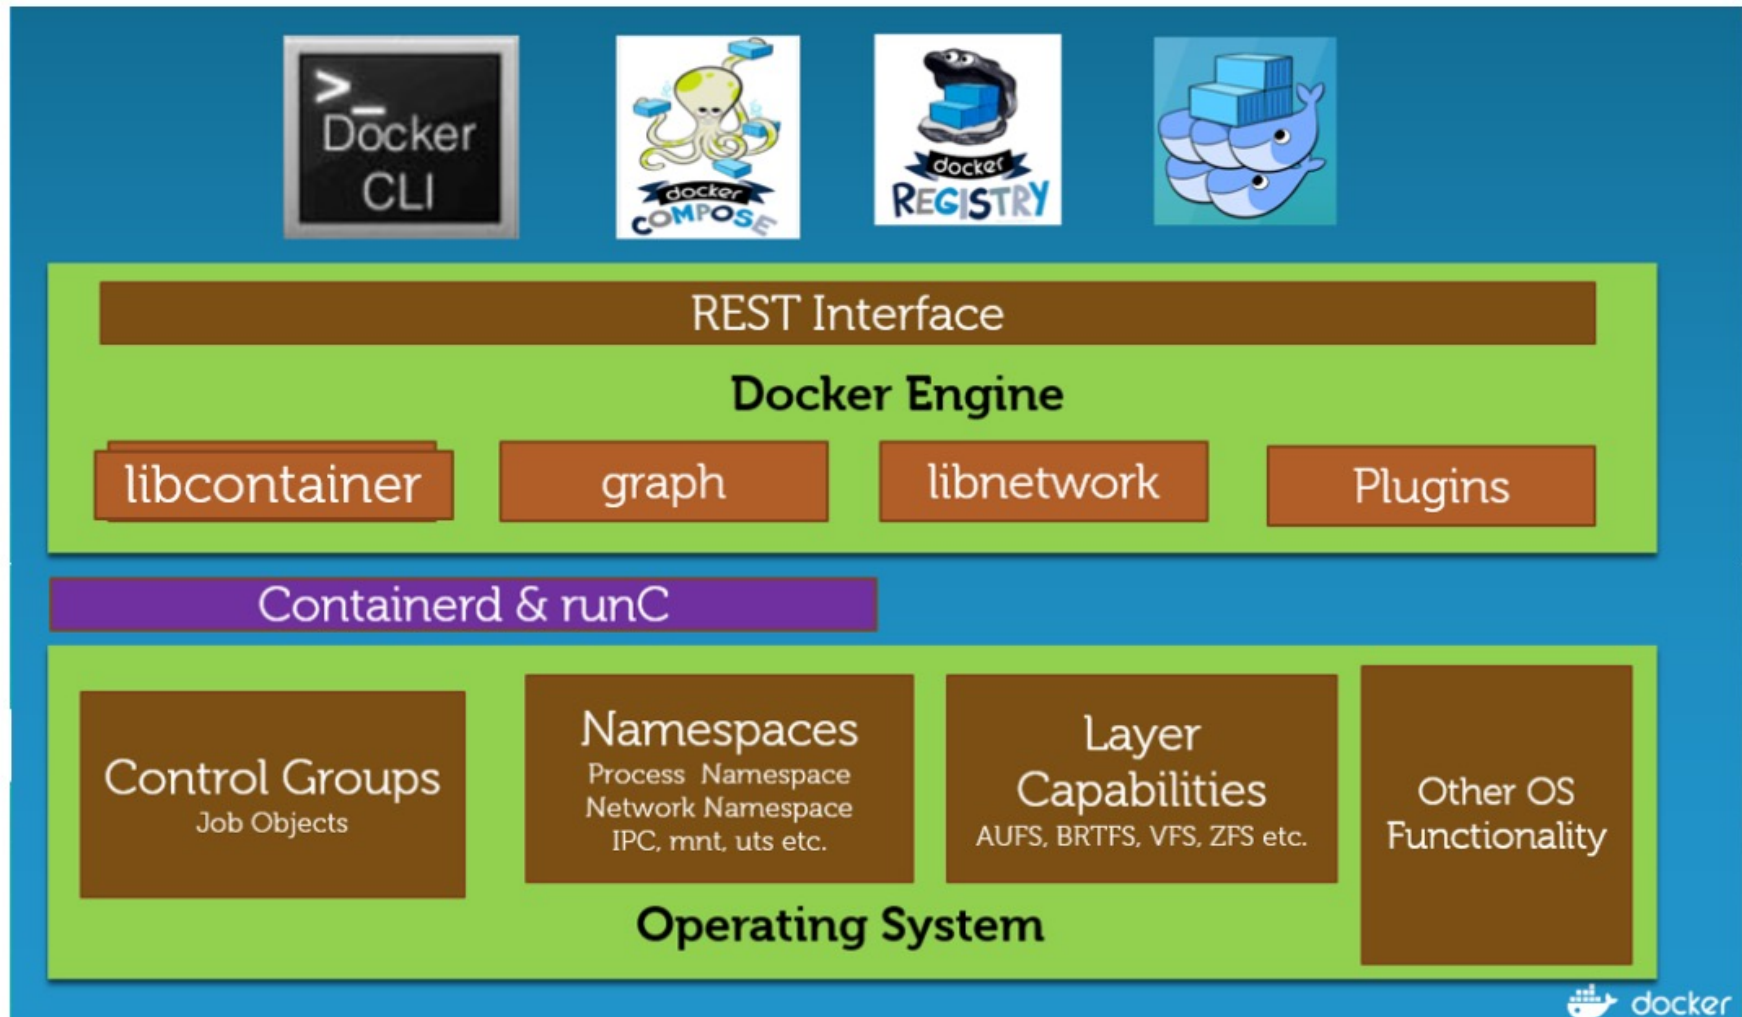
namespaces

cgroups

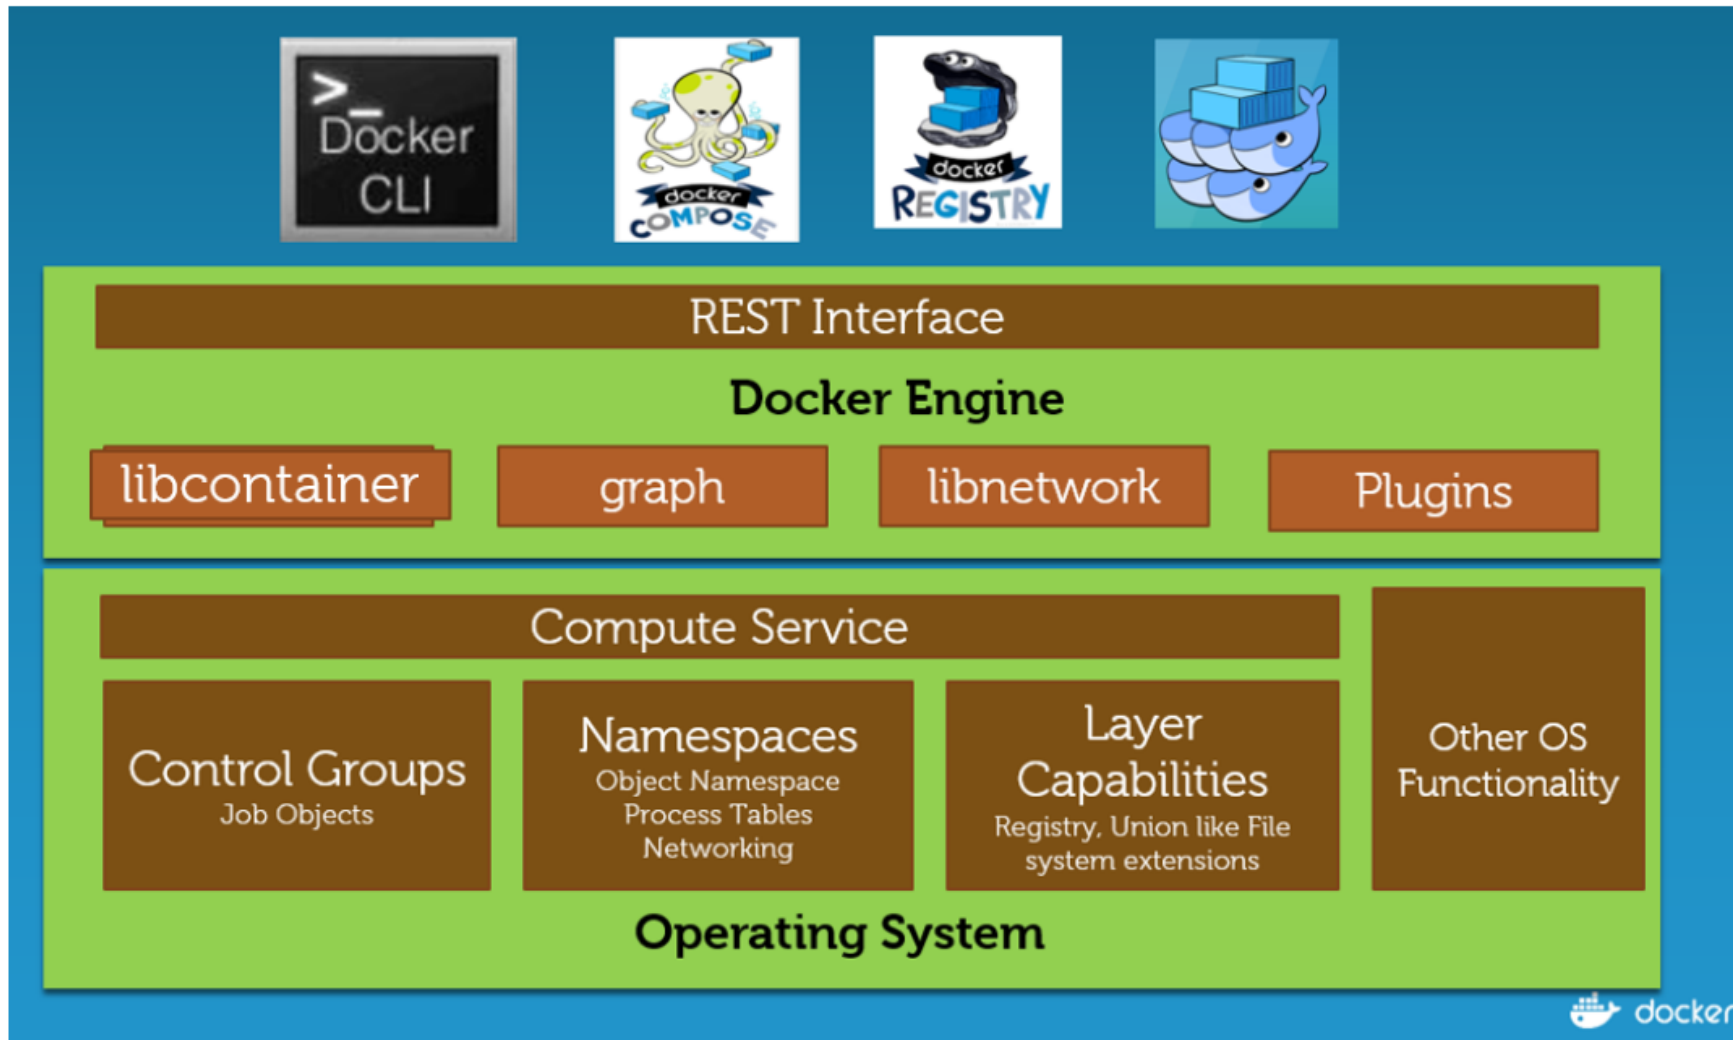
SELinux/AppArmor

Linux kernel

Docker Engine on Linux Platform



Docker Engine on Windows Platform



Docker Registry

- ▶ Central repository to store and deliver images
 - ▶ indexed by **name** and **tag**
 - ▶ can be private to an infrastructure or public (Docker Hub)
- ▶ Clients push images to the registry
 - ▶ Local client
- ▶ Docker Engine pulls images and executes the container
 - ▶ Running on servers
- ▶ Layers are hashed and indexed allowing reuse
 - ▶ Pulling an image only requires the layers that do not exist locally



Docker Hub - Largest Public Registry

Name

The screenshot shows the Docker Hub interface for the 'nginx' repository. A green arrow points from the 'Name' label to the 'nginx' repository name, which is highlighted with a red box. Another green arrow points from the 'Tags' label to the 'Supported tags and respective Dockerfile links' section, which is also highlighted with a red box. The interface includes a search bar, navigation links, and a list of supported tags.

dockerhub nginx Explore Sign In Pricing Sign Up

Explore nginx

nginx ☆
Docker Official Images
Official build of Nginx.

1B+

Container Linux 386 x86-64 ARM 64 IBM Z ARM PowerPC 64 LE Application Infrastructure

Official Image

Linux - x86 (latest)

Copy and paste to pull this image

```
docker pull nginx
```

[View Available Tags](#)

Description Reviews Tags

Supported tags and respective Dockerfile links

- 1.17.9, mainline, 1, 1.17, latest
- 1.17.9-perl, mainline-perl, 1-perl, 1.17-perl, perl
- 1.17.9-alpine, mainline-alpine, 1-alpine, 1.17-alpine, alpine

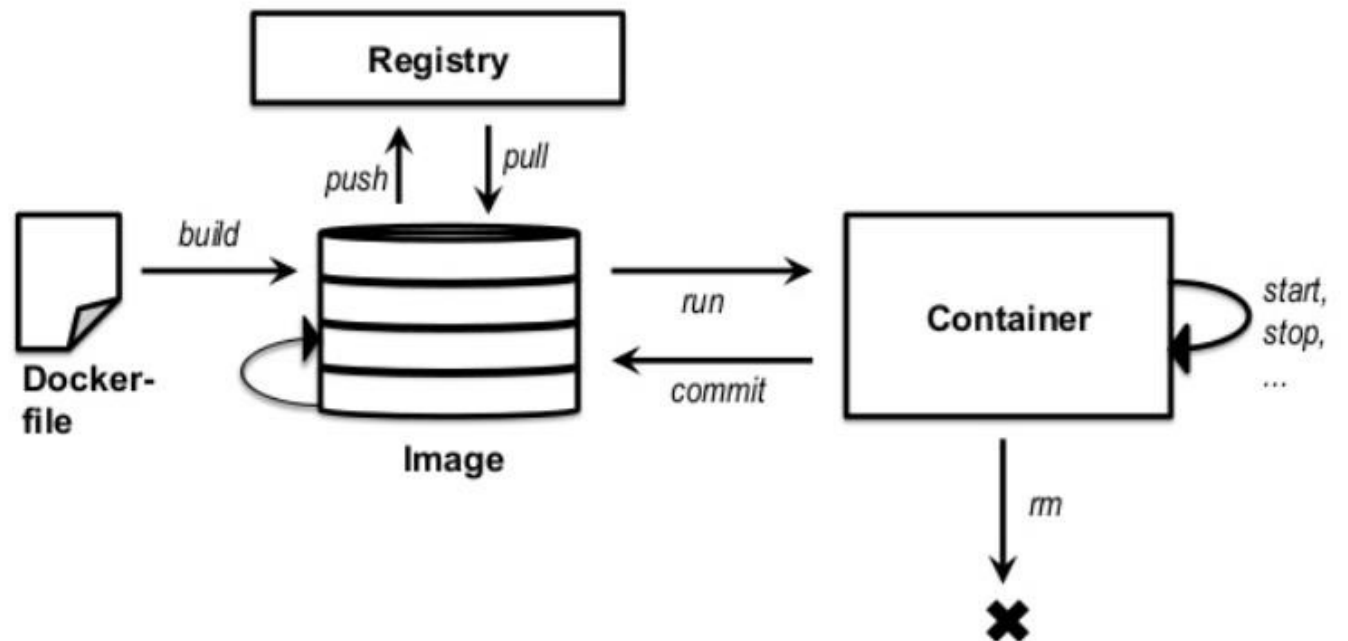
Tags

```
root@server:~# docker search nginx
```

| NAME | DESCRIPTION | STARS | OFFICIAL | AUTOMATED |
|----------------------------------|---|-------|----------|-----------|
| nginx | Official build of Nginx. | 12788 | [OK] | |
| jwilder/nginx-proxy | Automated Nginx reverse proxy for docker con... | 1750 | | [OK] |
| richarvey/nginx-php-fpm | Container running Nginx + PHP-FPM capable of... | 758 | | [OK] |
| linuxserver/nginx | An Nginx container, brought to you by LinuxS... | 95 | | |
| bitnami/nginx | Bitnami nginx Docker Image | 77 | | [OK] |
| tiangolo/nginx-rtmp | Docker image with Nginx using the nginx-rtmp... | 64 | | [OK] |
| jc21/nginx-proxy-manager | Docker container for managing Nginx proxy ho... | 45 | | |
| nginxdemos/hello | NGINX webserver that serves a simple page co... | 41 | | [OK] |
| nginx/unit | NGINX Unit is a dynamic web and application ... | 36 | | |
| jlesage/nginx-proxy-manager | Docker container for Nginx Proxy Manager | 35 | | [OK] |
| nginx/nginx-ingress | NGINX Ingress Controller for Kubernetes | 28 | | |
| privatebin/nginx-fpm-alpine | PrivateBin running on an Nginx, php-fpm & Al... | 22 | | [OK] |
| schmunk42/nginx-redirect | A very simple container to redirect HTTP tra... | 18 | | [OK] |
| blacklabelops/nginx | Dockerized Nginx Reverse Proxy Server. | 13 | | [OK] |
| nginxinc/nginx-unprivileged | Unprivileged NGINX Dockerfiles | 13 | | |
| centos/nginx-112-centos7 | Platform for running nginx 1.12 or building ... | 12 | | |
| raulr/nginx-wordpress | Nginx front-end for the official wordpress:f... | 12 | | [OK] |
| centos/nginx-18-centos7 | Platform for running nginx 1.8 or building n... | 12 | | |
| nginx/nginx-prometheus-exporter | NGINX Prometheus Exporter | 9 | | |
| sophos/nginx-vts-exporter | Simple server that scrapes Nginx vts stats a... | 7 | | [OK] |
| mailu/nginx | Mailu nginx frontend | 6 | | [OK] |
| bitnami/nginx-ingress-controller | Bitnami Docker Image for NGINX Ingress Contr... | 4 | | [OK] |
| ansibleplaybookbundle/nginx-apb | An APB to deploy NGINX | 1 | | [OK] |
| wodby/nginx | Generic nginx | 0 | | [OK] |
| centos/nginx-110-centos7 | Platform for running nginx 1.10 or building ... | 0 | | |

Workflow

- Search docker image
`$ docker search nginx`
- Pull docker image
`$ docker pull nginx`
- List local docker images
`$ docker image ls`
- Run container
`$ docker run -d --name frontend-server nginx`



Docker file

- ▶ **Sequence of commands** to prepare the container for execution
 - ▶ Lists dependencies
 - ▶ Lists commands to be executed inside container
 - ▶ Sets the command to execute on start
- ▶ Used locally or distributed in registry
 - ▶ Public registry: Docker hub
- ▶ May define versions for same software
 - ▶ nginx:latest, nginx:perl, nginx:alpine...

```
FROM debian:stretch-slim

LABEL maintainer="NGINX Docker Maintainers <docker-maint@nginx.com>"

ENV NGINX_VERSION 1.15.9-1~stretch
ENV NJS_VERSION 1.15.9.0.2.8-1~stretch

RUN set -x \
    && apt-get update \
    && apt-get install --no-install-recommends --no-install-
suggests -y gnupg1 apt-transport-https ca-certificates \
    && \

...OMMITED...

RUN ln -sf /dev/stdout /var/log/nginx/access.log \
    && ln -sf /dev/stderr /var/log/nginx/error.log

EXPOSE 80

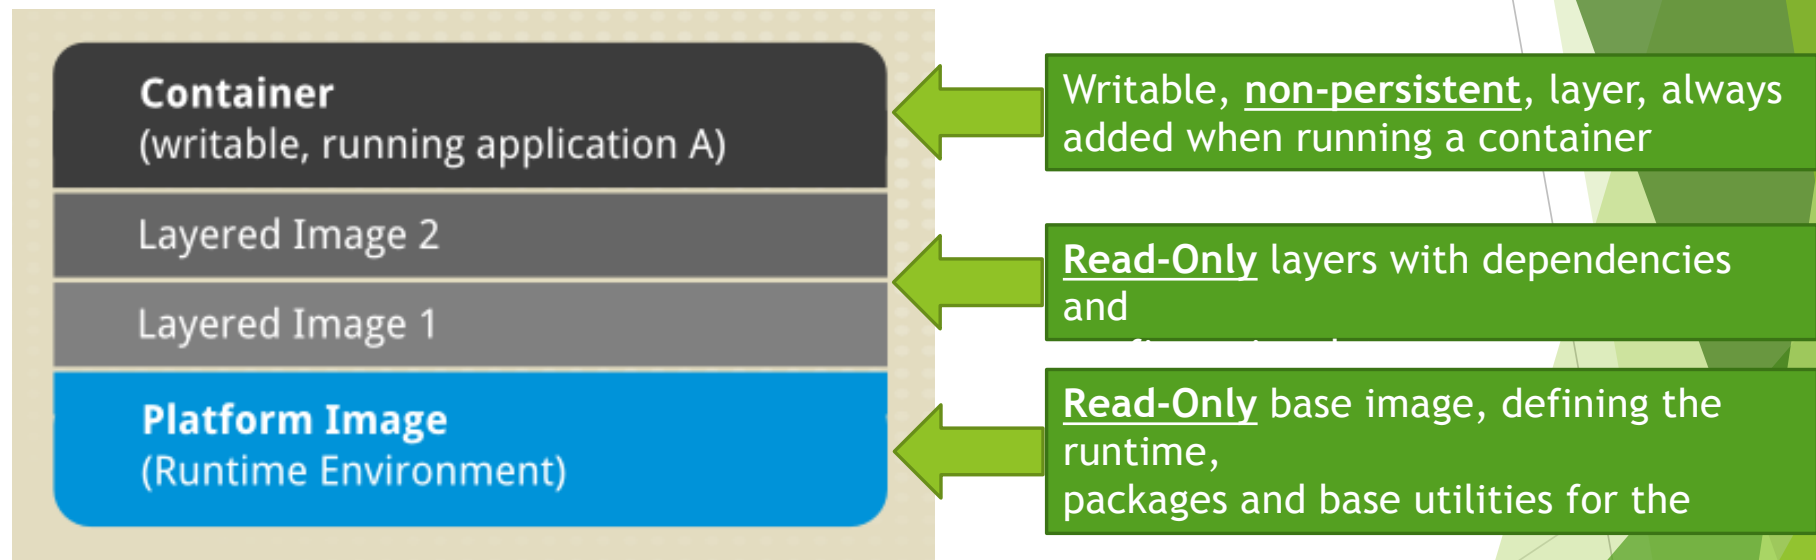
STOPSIGNAL SIGTERM

CMD ["nginx", "-g", "daemon off;"]
```

Images

- ▶ Containers have their initial data in images
 - ▶ Contain data, applications, libraries
 - ▶ Contain an entry point to be executed when the container is initiated
- ▶ Composed by multiple layers
 - ▶ A base image
 - ▶ Several images, changing the content (with differences)
 - ▶ Exploits overlayfs (a unionfs similar to AUFS)
 - ▶ Can use others like btrfs, ZFS
- ▶ Read Only and Non persistent
 - ▶ Serve as templates for containers to start
 - ▶ Multiple containers will use the same images (containers are ephemeral)

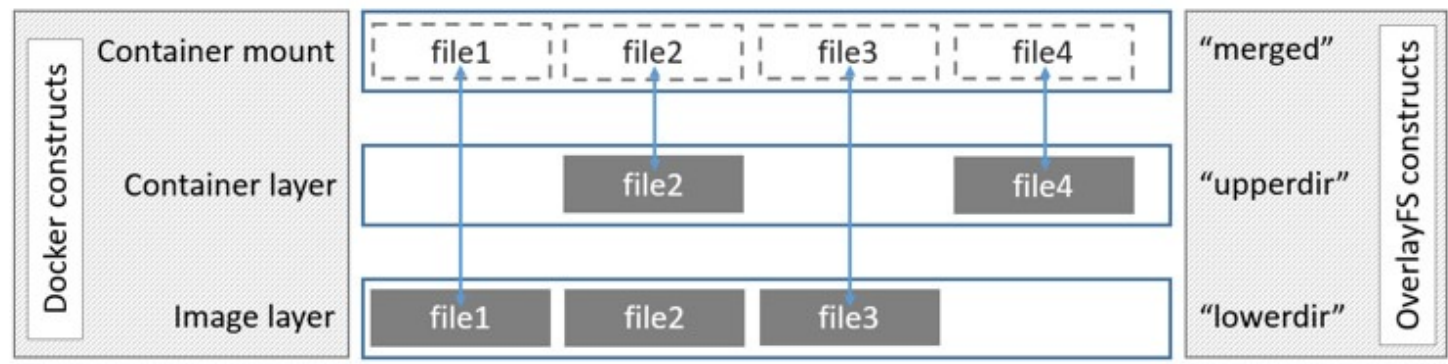
Image layering



Available at `/var/lib/docker/...`

Images

- ▶ It's important to keep images small
 - ▶ **Dont:** base your image on a full blown distro (e.g. debian)
 - ▶ **Do:** base your image on a small distro (e.g. Alpine)
- ▶ Presented on the host
 - ▶ `/var/lib/docker/overlay2/<ID>/`
 - ▶ `/merged`: FS seem inside container
 - ▶ `/diff`: differences from base



Images


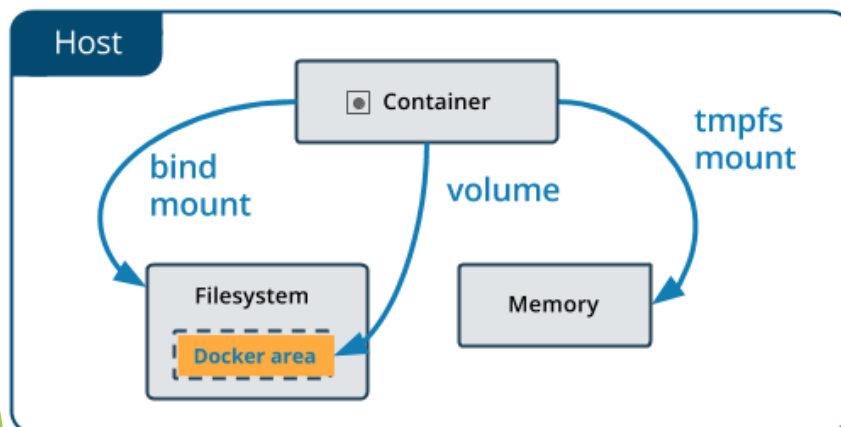
- ▶ OverlayFS has limitations!
 - ▶ Only a subset of the calls is supported
 - ▶ May create issues with some applications or decrease performance

▶ Issues:

- ▶ open: opening a file RW after the same file is open as R will fail.
 - ▶ first open (R) will refer to a file in the base image
 - ▶ second open (RW) will result in a copy of the file being created (stored in the diff folder)
 - ▶ File descriptions will not point to the same file
- ▶ rename: not supported
 - ▶ applications must copy and then unlink
 - ▶ much slower and will take additional disk space

Container Persistence

- ▶ Top layer is writable, but not persisted!
 - ▶ Stopping the container will result in data loss
- ▶ Persistence must be achieved through external resources
 - ▶ Bind mount: Host path
 - ▶ Volume: Docker managed resource
 - ▶ ... or external, database cluster



Container Persistence

- ▶ Bind mounts:
 - ▶ mount an existing path into a path inside the container
- ▶ Example: mount *app* dir from host, into */var/www* inside the container
 - ▶ Usefulness: persistence or add shared/specific data into container (ex, web page, static resources)

```
$ docker run -d\  
  --mount type=bind,source=/app,target=/var/www\  
  -p 8080:80  
  --name frontend-server  
  nginx
```

Container Persistence

- ▶ Volume mounts:
 - ▶ mount a specific storage object (managed by docker) at a given path
 - ▶ volumes are not deleted with containers
- ▶ Example: mount *app* volume, into */var/www* inside the container as read-only
 - ▶ Usefulness: persistence or add shared/specific data into container (ex, web page, static resources)

```
$ docker run -d\  
  --mount type=volume,source=app:/var/www:ro\  
  -p 8080:80  
  --name frontend-server  
  nginx
```


Container Network

- ▶ Containers run with specific network configuration
 - ▶ **No ingress network access added by default**
- ▶ Standard physical networks (others can be created)
 - ▶ bridge: virtual switch
 - ▶ host: network only connected to host OS
 - ▶ none: isolated network (only with loopback, and can be customized)
- ▶ Others
 - ▶ **overlay: used to provide connectivity over multiple engines (swarm mode)**
 - ▶ **macvlan: assigns specific MAC address to container**

Container Network Ports

- ▶ Services inside container cannot provide services to other applications
 - ▶ No network services by default
- ▶ Run command must expose ports
 - ▶ Outside ports mapped to inside ports
- ▶ E.g. map external port 8000 to internal (inside container) port 80

```
$ docker run -d -p 8000:80 --name frontend-server nginx
```

Sensitive data

- ▶ Frequently, there is sensitive data that must be available to containers
 - ▶ Database/ext service access credentials, certificates, admin credentials
 - ▶ Sensitive data is container specific, or shareable among several containers
- ▶ Direct (**no so correct**) approach: Build images with credentials
 - ▶ must build different images for each credential pair
 - ▶ must rebuild images when credentials change
 - ▶ credentials frequently end in the teams source control (e.g. Git repository)

Docker Compose



- ▶ Compose services created from multiple containers
 - ▶ that is... consider the specification of a complete service, and not
 - ▶ docker-compose.yml file
- ▶ Example:
 - ▶ specifies service name: nginx
 - ▶ container name and build ref
 - ▶ volumes to add
 - ▶ uses an environmental variable
 - ▶ ports to expose to outside

```
version: '3'
services:
  nginx:
    container_name: nginx
    build:
      context: ./nginx
      dockerfile: Dockerfile
    volumes:
      - app:/var/www
      - ${NGINX_LOG_PATH}:/var/log/nginx
    ports:
      - 80:80
      - 443:443
```

Sensitive data: Docker Secrets

- ▶ Allows creating blobs of data with restricted access
 - ▶ Exist independently of the containers
 - ▶ Allow changing the secret without changing the container
 - ▶ ex: database access creds for dev, staging and prod
 - ▶ Only accessed inside specific containers
 - ▶ at `/run/secrets/secret_name`
- ▶ Manage secret: `docker secret create/inspect/list/rm`
- ▶ Add to service: `docker service... --secret secret_name`
 - ▶ Will expose `secret_name` inside container

Docker Compose



```
version: '3'
services:
  nginx:
    container_name: nginx
    build:
      context: ./nginx
      dockerfile: Dockerfile
    ...
  secrets:
    key_file:
      file: server-key.pem
    cert_file:
      file: server-cert.pem
```

Kubernetes



Kubernetes



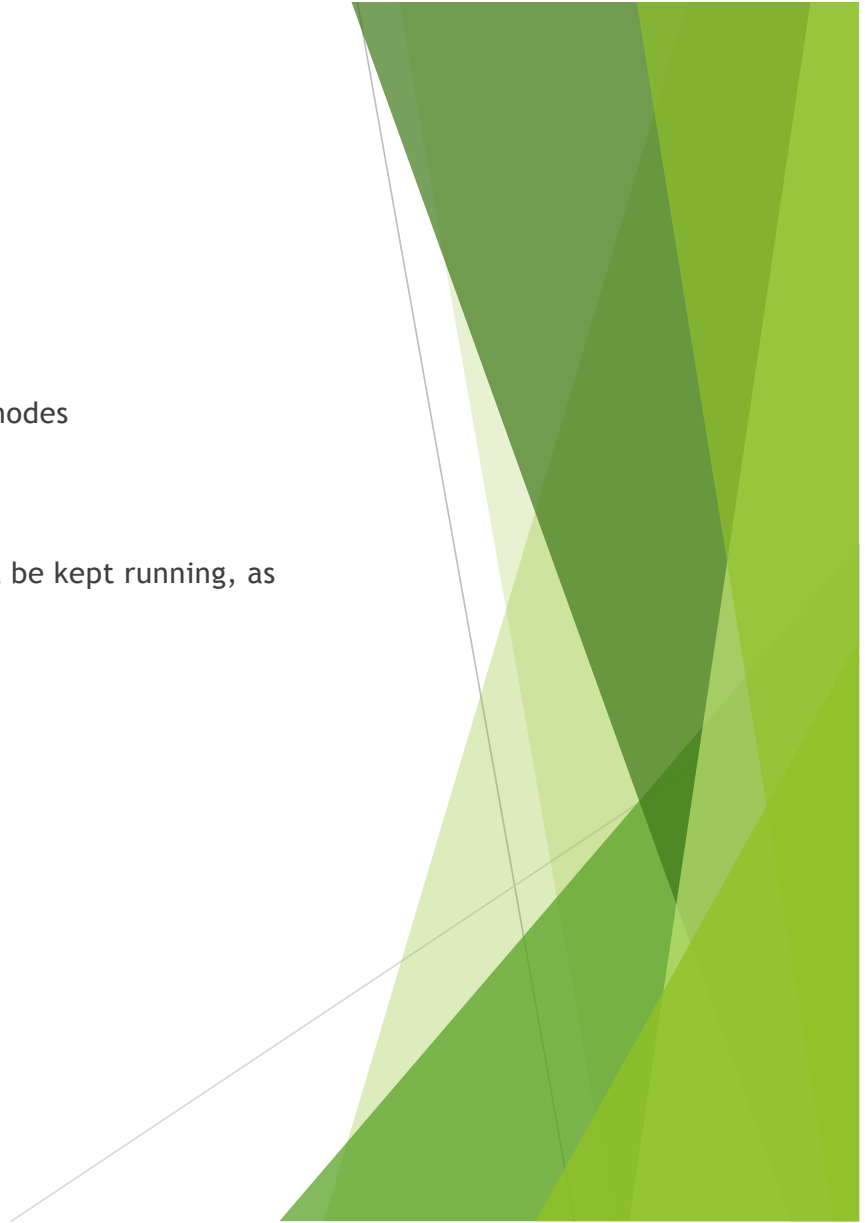
- ▶ A platform for **orchestrating (Docker/Containerd) containers** in a **clustered environment**, usually with **multiple hosts**
- ▶ Provides:
 - ▶ **container grouping**: groups containers in same environment
 - ▶ **load balancing** distributes load across physical resources
 - ▶ **auto-healing**: recover from failures (auto start containers)
 - ▶ **scaling features**: Increase number of containers as required by load
 - ▶ ...
- ▶ **Motivation**: running individual containers is not enough. Kubernetes is an **orchestration solution that keeps services operating**

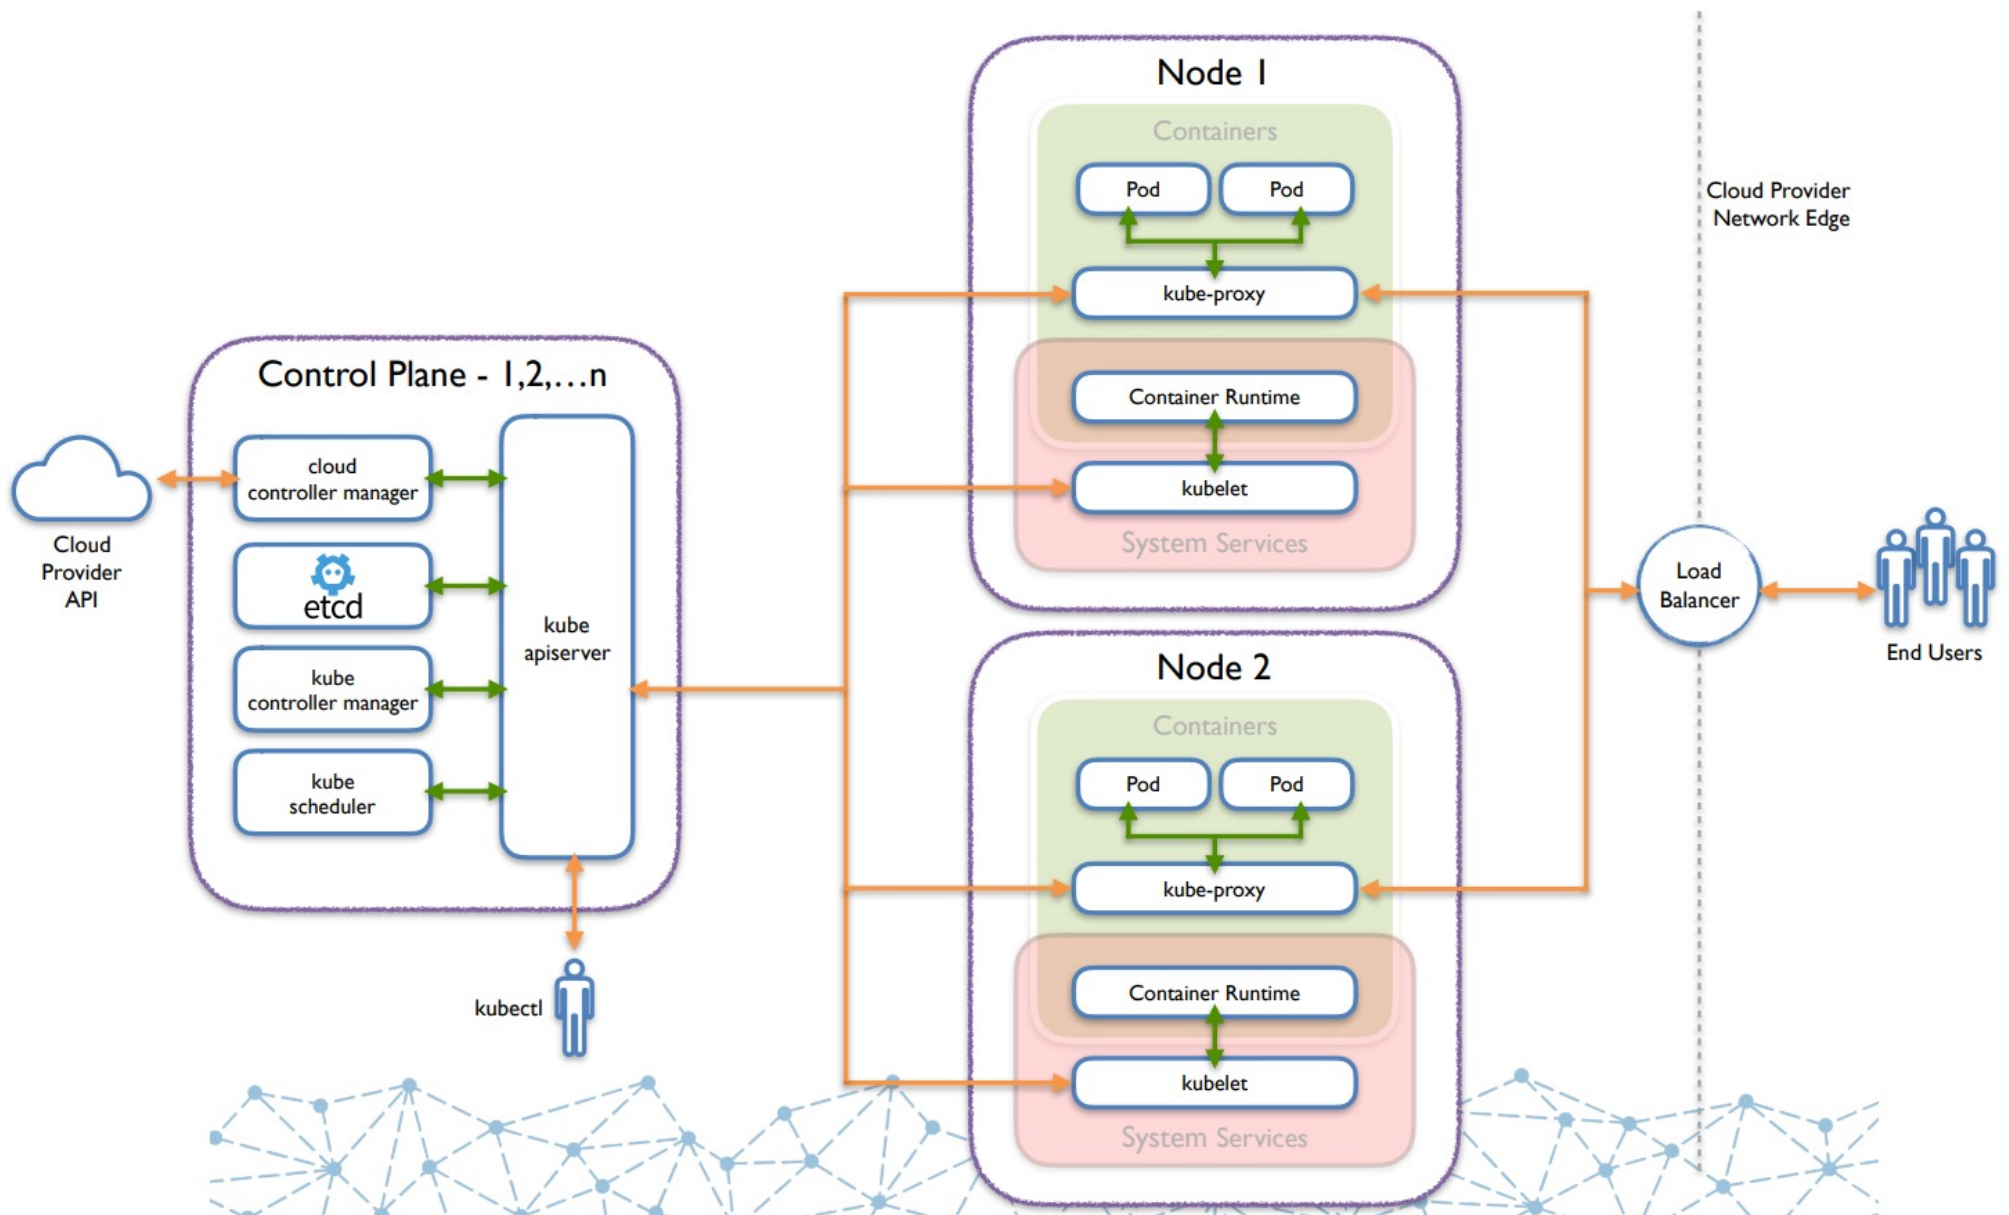
Kubernetes

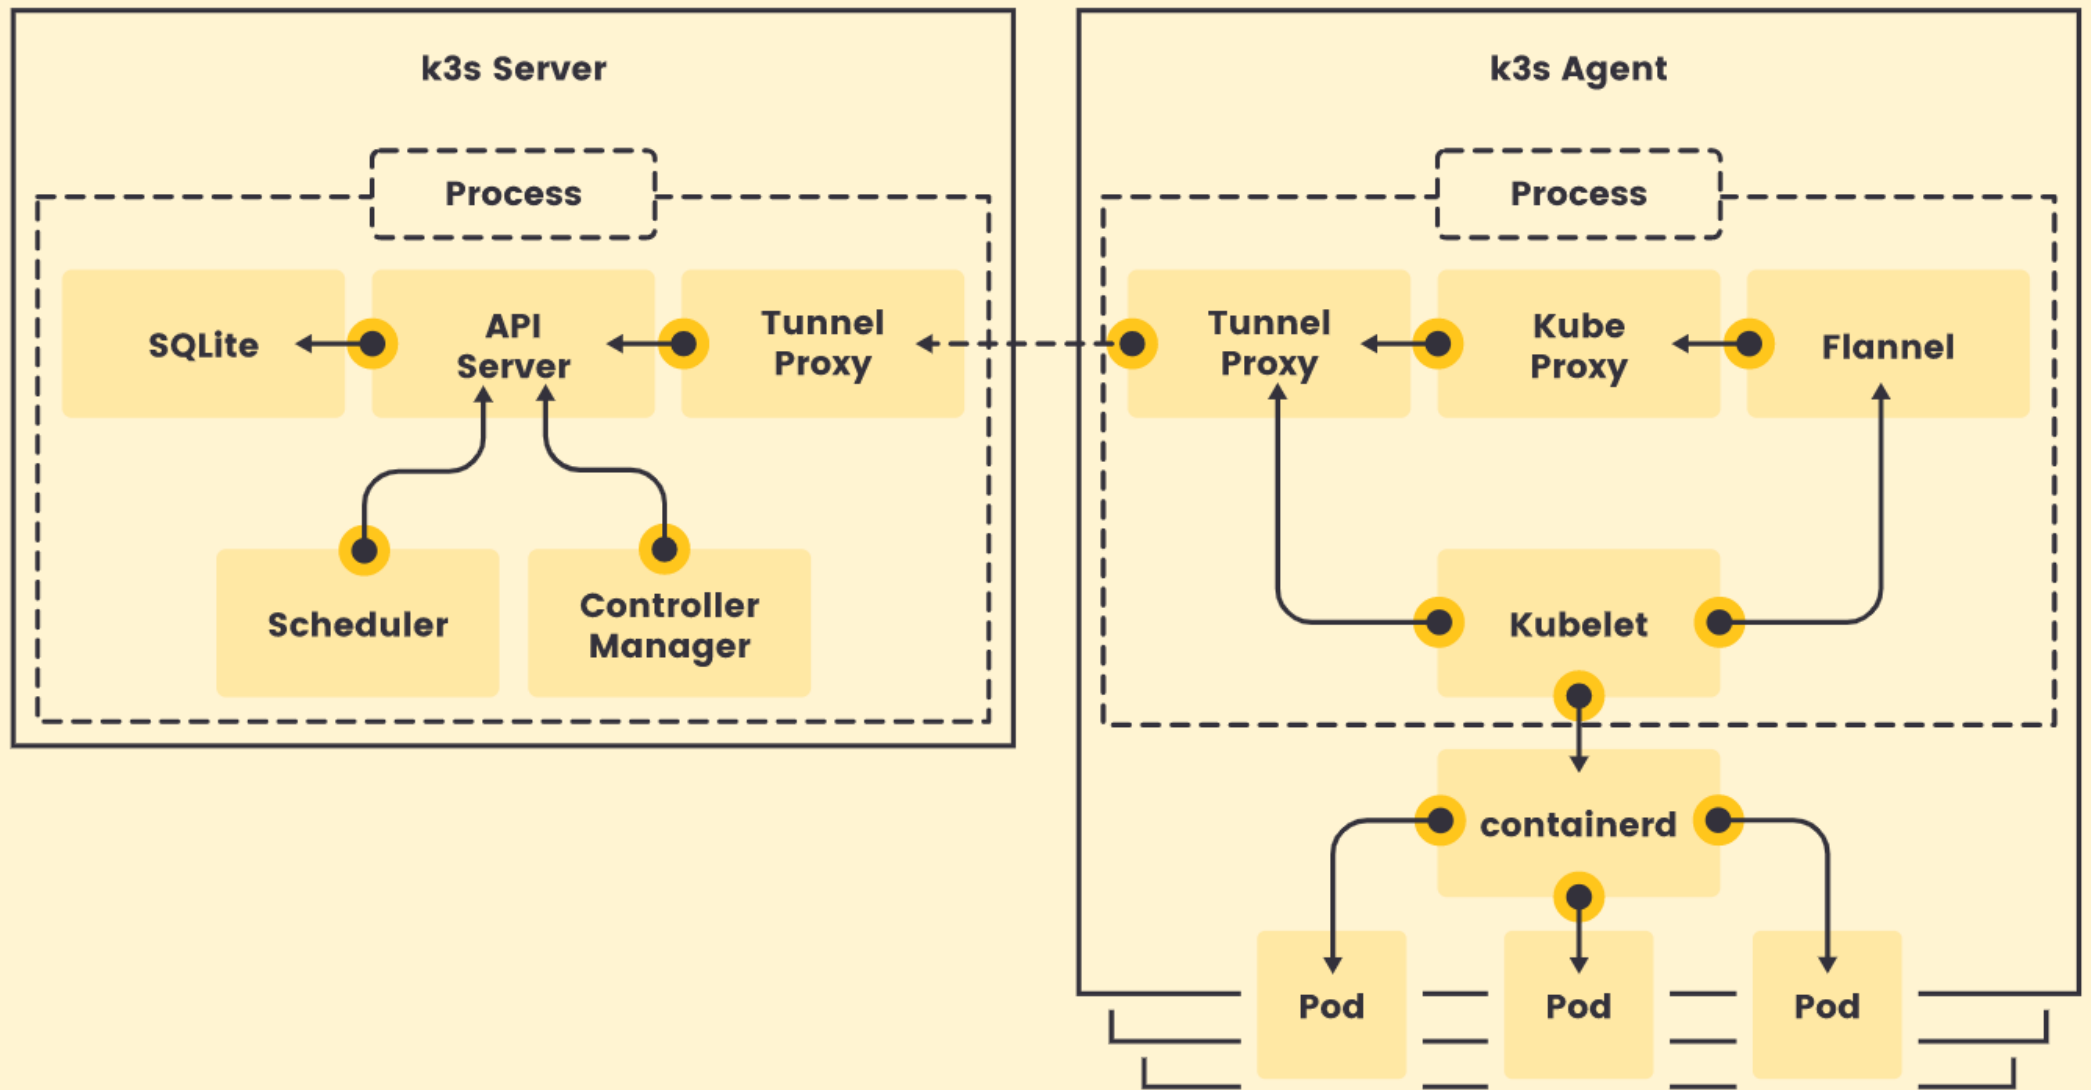
- ▶ An orchestrator manages the lifecycle of an entity (VM, container, etc..)
 - ▶ In this case: the lifecycle of containers
 - ▶ How it is created
 - ▶ Where it is instantiated
 - ▶ How it is balanced to new nodes
- ▶ Orchestrators usually also manage related resources
 - ▶ Data volumes
 - ▶ Secrets
 - ▶ Configuration values
 - ▶ External resources and services
- ▶ Orchestrators are not simple automatic deployment tools
 - ▶ They have an integrated view of the resources and make decisions to optimize placement

Kubernetes

- ▶ Focus on managing hundreds/thousands of containers, on a fleet of thousands of nodes
 - ▶ While dealing with rollout of new versions, failure, security, ...
- ▶ Implement a software defined platform: given the specification, the software will be kept running, as long as there are resources







Key software components

- ▶ Kube-apiserver: provides a REST API for clients to interact with
 - ▶ Control plane
 - ▶ Data Storage
- ▶ Etcd: The Cluster Datastore, with a consistent and highly available key value storage
 - ▶ For configuration data
- ▶ Kube-Controller-Manager: Manages all component control loops
 - ▶ Ensures the cluster is at the desired state
- ▶ Kube-Scheduler: Places new Pods into Nodes
 - ▶ Takes in consideration MANY scheduling decisions (affinity, anti-affinity, constraints, data locality....)

Key software components

- ▶ Kube-Proxy: Manages network rules at each node, and does connection forwarding / load balancing
 - ▶ Similar to the docker proxy
- ▶ Kubelet: takes in consideration the description of each Pod (PodSpecs) and ensures (locally in the node) that the containers are running as expected
- ▶ Container Runtime Engine: A software that runs containers, many supported
 - ▶ Containerd (docker)
 - ▶ Rkt
 - ▶ cri-o
 - ▶ Virtlet
 - ▶ Kata

Networking


- ▶ Pods have an internal network namespace
 - ▶ You have multiple pods with the same networks!
- ▶ Pod Network: Cluster wide network for Pod-to-Pod communication
- ▶ Service Network: Virtual IPs (durable), managed by kube-proxy for service discovery
 - ▶ Because Pod IP address are volatile, services provide anchor points for communication

Namespaces

- ▶ Logical Clusters inside the Real Cluster
- ▶ Provide the means for partitioning the cluster to multiple teams/projects
- ▶ Objects are created under a namespace and deletion of the namespace may also delete all objects
 - ▶ Exceptions are the resources related StatefulSets (later)
- ▶ Some namespaces are always present (default, kube-system,...)
 - ▶ Typical namespaces: production, development, monitoring

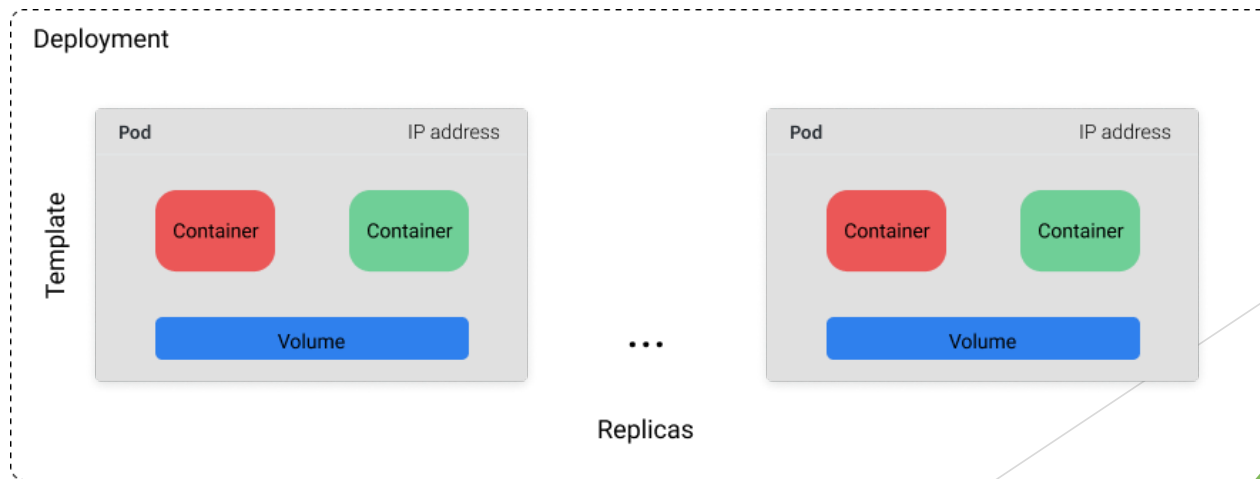
Pod

- ▶ A group of Containers. A basic management unit.
 - ▶ Includes information about what to run, and how to run something
 - ▶ Mimics a local host with applications (containers)
- ▶ Containers are ephemeral
- ▶ Containers are tagged to versions and we go changing version as the infrastructure evolves
- ▶ Pod contents are always co-located and co-scheduled



Deployment

- ▶ A set of pods that should operate together
 - ▶ Can be a software, but also includes volumes, networks, secrets, etc..
- ▶ Defined using a declarative format (yaml file) stating the desired state of the pods



Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app
  namespace: diving-into-k3s
spec:
  replicas: 1
  selector:
    matchLabels:
      app: app
  template:
    metadata:
      labels:
        app: app
    spec:
      containers:
        - name: app
          image: 10.110.0.3:5000/app:v1
          resources:
            requests:
              memory: "32Mi"
              cpu: "10m"
            limits:
              memory: "128Mi"
              cpu: "500m"
          ports:
            - containerPort: 8080
```

What is the context

How many instances

How to find pods in an infrastructure

What to run. In this case, a Docker container

Resource limits


How to expose

StatefulSet

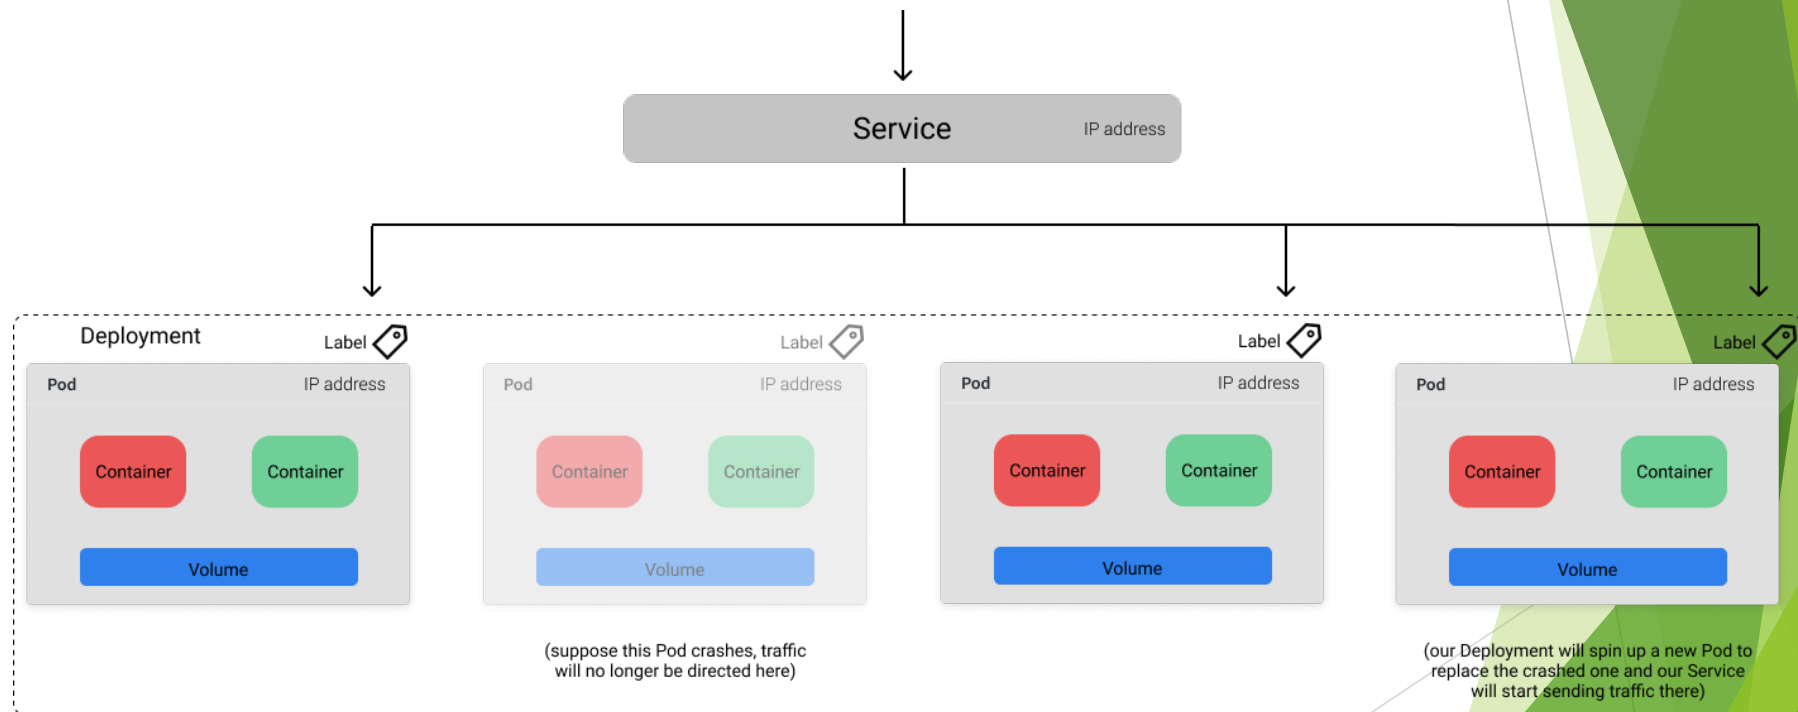
- ▶ Similar to a deployment but Pods are stateful
 - ▶ Have persistent and unique IDs
 - ▶ Are started by a specific order
 - ▶ Storage is persistent
 - ▶ Requires a StorageClass to actually store data
 - ▶ Graceful deployment and scaling
- ▶ Useful for stateful software, whose instances have unique roles
 - ▶ e.g a Database

Service

- ▶ K8s abstracts internal aspects from users, such as the IP addresses of replicas
 - ▶ Also, number and location of said replicas
- ▶ A service provides the visibility to reach a deployment **from inside the cluster**
 - ▶ Allows Pods to reach other pods
 - ▶ Service provides a stable endpoint, no matter how many pods are available (or where)
- ▶ Persistent
 - ▶ Static cluster IP
 - ▶ Static DNS Name



Service



Service

- ▶ A service named “app” allows Pods to access other pods by DNS, through the service
 - ▶ e.g. curl <http://app> would be redirected to one pod (round robin)

```
apiVersion: v1
kind: Service
metadata:
  name: app
  namespace: diving-into-k3s
spec:
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: app
```

Name and context

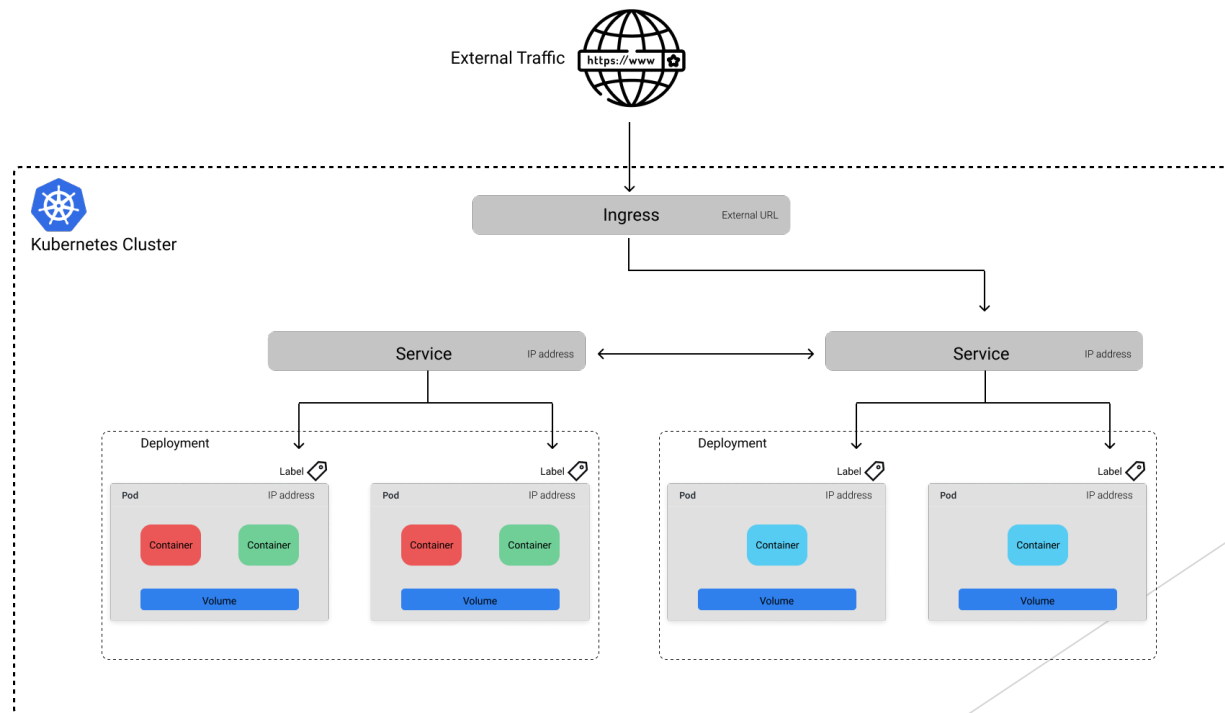
Which port is exposed

Which port used to reach Pods

Which pods are used

Ingress

- ▶ Similar to a Service, makes some services visible to the outside world
- ▶ Frequently uses a proxy software (Nginx, Traefik)



Ingress with Traefik

Name and Namespace

Traefik Specific Configuration for This ingress

```
---
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: app-k3s
  namespace: diving-into-k3s
  annotations:
    kubernetes.io/ingress.class: traefik
    traefik.ingress.kubernetes.io/frontend-entry-points:
      http,https
    traefik.ingress.kubernetes.io/redirect-entry-point:
      https
    traefik.ingress.kubernetes.io/redirect-permanent:
      "true"
spec:
  rules:
  - host: app.k3s
    http:
      paths:
      - path: /
        backend:
```

How to expose the Pods (the hostname and port)

What Pods to use (will use round robin)

Job

- ▶ An ephemeral execution to run one time
- ▶ In contrast to a pod in a deployment, which are expected to be running, jobs are one-time tasks.
- ▶ Kubernetes ensures the job runs as stated, dealing with availability of

```
apiVersion: batch/v1
kind: Job
metadata:
  name: model-train-job
spec:
  template:
    spec:
      containers:
      - name: ml-model-train
        image: ml-development:1.2.1
        command: ["python", "train.py"]
        restartPolicy: Never
      backoffLimit: 5
```

How should we create the Pod for this Job?

How many times should we try to complete this Job before quitting?

ConfigMaps

- ▶ Provide the means to create generally available configurations for your entities
 - ▶ Used for non-confidential data
 - ▶ Key-Value pairs
 - ▶ Can be injected using configuration files in volumes, command line arguments or environmental variables
- ▶ Allows decoupling the software from the execution environment specific configuration
 - ▶ Usually, different environments (dev and prod?) may have different configurations
- ▶ ConfigMaps are created, and then attributed to a Pod in the deployment
 - ▶ Only the required Pods will access the configuration

ConfigMaps

A subPath (the Item) in a Volume is mapped to the Containers.

The config file (nginx.conf) will be available at /etc/nginx/nginx.conf

ConfigMap Items are mapped to a Volume.

Definition of the ConfigMap (notice that it is restricted to the namespace)

ConfigMap has a name, and items inside it

```
# Create nginx deployment
---
...
  containers:
    - image: nginx:alpine
      name: nginx
      ports:
        - containerPort: 80
      resources: {}
      volumeMounts:
        - name: nginx-conf
          mountPath: /etc/nginx/nginx.conf
          subPath: nginx.conf
          readOnly: true
  volumes:
    - name: nginx-conf
      configMap:
        name: nginx-conf
        items:
          - key: nginx.conf
            path: nginx.conf
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-conf
  namespace: diving-into-k3s
data:
  nginx.conf: |
    ...content of the nginx.conf configuration
```

Secrets

- ▶ Similar to ConfigMaps, but aims to store a small, sensitive information blob (password, token, key)
 - ▶ By default are stored as a base64 string, with access limited only by the API
 - ▶ Kubernetes allows both Encryption and Role Based Access Control (RBAC)
 - ▶ Secrets will be strongly encrypted and only provided to allowed users
 - ▶ Differ from ConfigMaps as:
 - ▶ They can only be used by kubelet, an environmental variable or a volume
 - ▶ Cannot pass it through arguments
 - ▶ They have a type, defining its purpose

Secrets

| Builtin Type | Usage |
|-------------------------------------|---------------------------------------|
| Opaque | arbitrary user-defined data |
| kubernetes.io/service-account-token | service account token |
| kubernetes.io/dockercfg | serialized ~/.dockercfg file |
| kubernetes.io/dockerconfigjson | serialized ~/.docker/config.json file |
| kubernetes.io/basic-auth | credentials for basic authentication |
| kubernetes.io/ssh-auth | credentials for SSH authentication |
| kubernetes.io/tls | data for a TLS client or server |
| bootstrap.kubernetes.io/token | bootstrap token data |

- ▶ The Opaque type has no restrictions on content
 - ▶ Meant for arbitrary user data loaded from a file, or specified interactively
- ▶ The remaining types impose specific content format